

Time Series Analysis

ARMA Models

Nicoleta Serban, Ph.D.

Professor

Stewart School of Industrial and Systems Engineering

ARMA Modeling: Data Example

About This Lesson



Emergency Department Care

Have you ever experienced long waits in the Emergency Department?

- Good predictions of daily inflow in an emergency department can assist in staffing and diversion
- Time series modeling can be useful in achieving good predictions.



Case Study Overview

Objective:

- Identify temporal patterns in the Emergency Department (ED) volume of patients
- Develop a model to predict ED volume

Time Series Data:

- Daily number of patients visiting an emergency department of a hospital in the Atlanta area with observations from 2010 until mid 2015
- Other predicting variables were made available by the hospital but we will only focus on the predictability of the time series with respect to temporal factors


Model Trend and Seasonality

Model Trend + Monthly Seasonality

```
library(mgcv)
time.pts = c(1:length(Volume))
time.pts = c(time.pts - min(time.pts))/max(time.pts)
month = as.factor(format(dates,"%b"))
week = as.factor(weekdays(dates))
gam.fit.seastr = gam(Volume.tr~s(time.pts)+month+week)
vol.fit.gam.seastr = fitted(gam.fit.seastr)
resid.process = Volume.tr-vol.fit.gam.seastr

acf(resid.process,lag.max=12*4,main="ACF: Residual Plot")
pacf(resid.process,lag.max=12*4,main="PACF: Residual Plot")
```

Model Trend and Seasonality



Fitting an AR process

Fit an AR(p) process for for $p \leq \text{order.max}$

`mod = ar(resid.process, order.max=20)`

What is the selected order?

`print(mod$order)`


Find the list of arguments provided by AR fit

`summary(mod)`

Plot AIC values on the log scale to easily identify minimum

`plot(c(0:20), mod$aic+.0001, type="b", log="y", xlab="order", ylab="log-AIC values")`

Fitting an AR process



Fitted AR process: Properties

Are the roots of fitted AR within the unit circle?

Extract roots from the model output:

```
roots = polyroot(c(1,(-mod$ar)))
```


Adjust the x and y -axis limits to include the full circle

```
plot(roots,xlim=c(-1.5,1.5),ylim=c(-1.5,1.5))
```

Draw a unit root circle

```
lines(complex(arg = seq(0,2*pi,len=300)))
```

Fitted AR process: Properties



Fitted AR process: Residual Analysis

Obtain the standardized residuals

```
resids = mod$resid[(mod$order+1): length(mod$resid)]
```

Plot the residuals

```
par(mfrow=c(2,2))
```

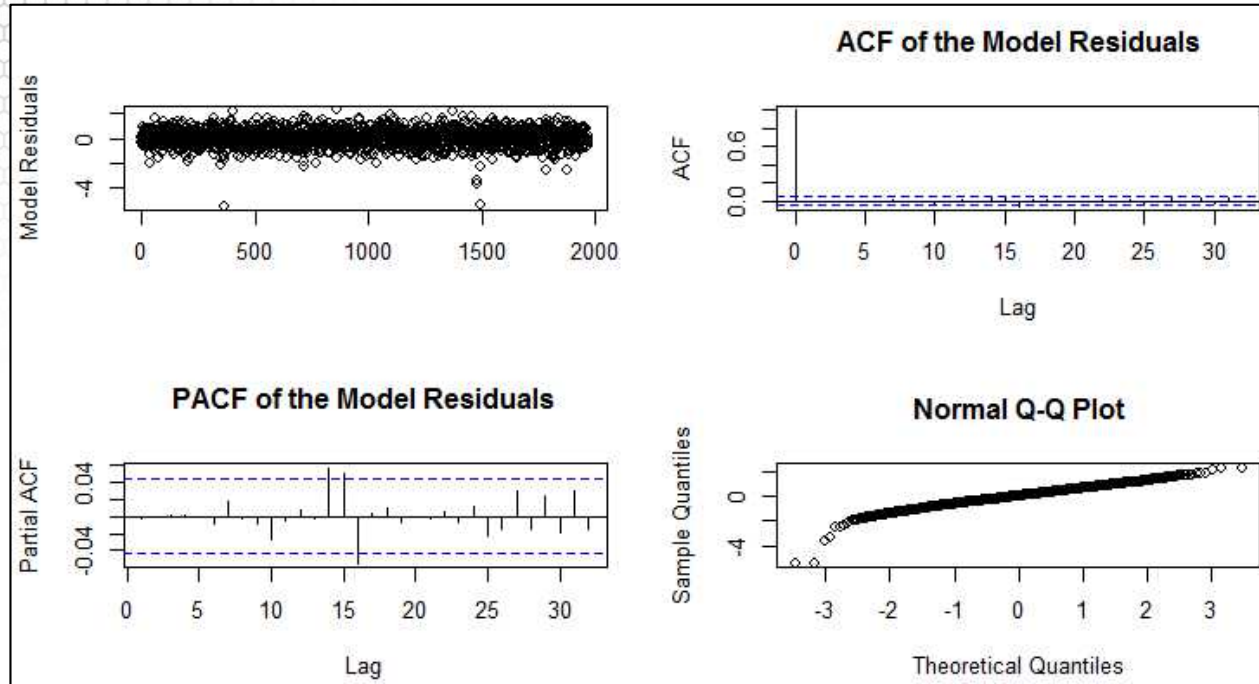
```
plot(resids,xlab="",ylab="Model Residuals")
```

```
acf(resids,main='ACF of the Model Residuals')
```

```
pacf(resids,main='PACF of the Model Residuals')
```

```
qqnorm(resids)
```

Fitted AR process: Residual Analysis



Fit an ARMA Model

Fit ARMA(6,1)

```
modarma = arima(resid.process, order = c(6,0,1),method = "ML")
```

Residual Analysis

```
par (mfrow=c(2,2))
```

```
plot(resid(modarma), ylab='Standardized Residuals')
```

```
abline(h=0)
```

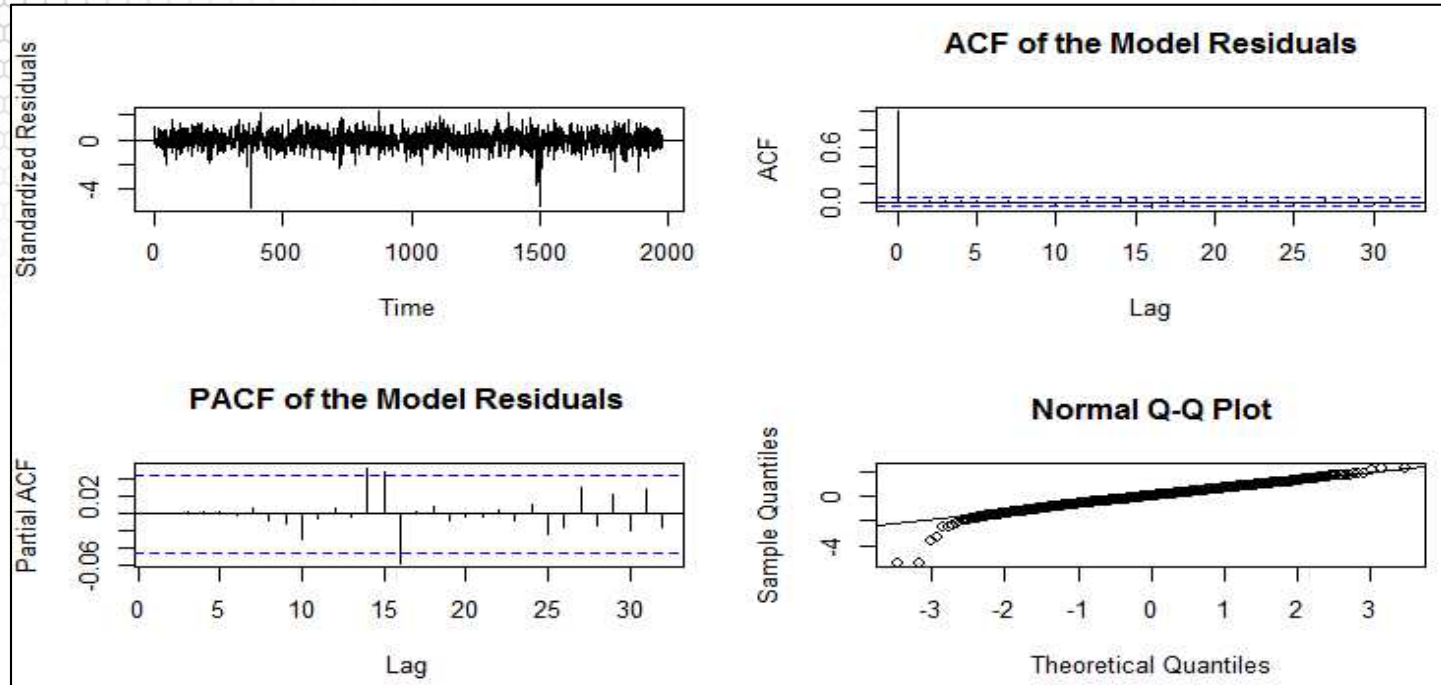
```
acf(as.vector(resid(modarma)),main= 'ACF of the Model Residuals')
```

```
pacf(as.vector(resid(modarma)),main='PACF of the Model Residuals')
```

```
qqnorm(resid(modarma))
```

```
qqline(resid(modarma))
```

Fit an ARMA Model



ARMA Model: Order Selection

Order selection – AIC

```
n = length(resid.process)
norder = 6
p = c(1:norder)-1; q = c(1:norder)-1
aic = matrix(0,norder,norder)
for(i in 1:norder){
  for(j in 1:norder){
    modij = arima(resid.process,order = c(p[i],0,q[j]), method='ML')
    aic[i,j] = modij$aic-2*(p[i]+q[j]+1)+2*(p[i]+q[j]+1)*n/(n-p[i]-q[j]-2)
  }
}
```


ARMA Model: Order Selection (cont'd)

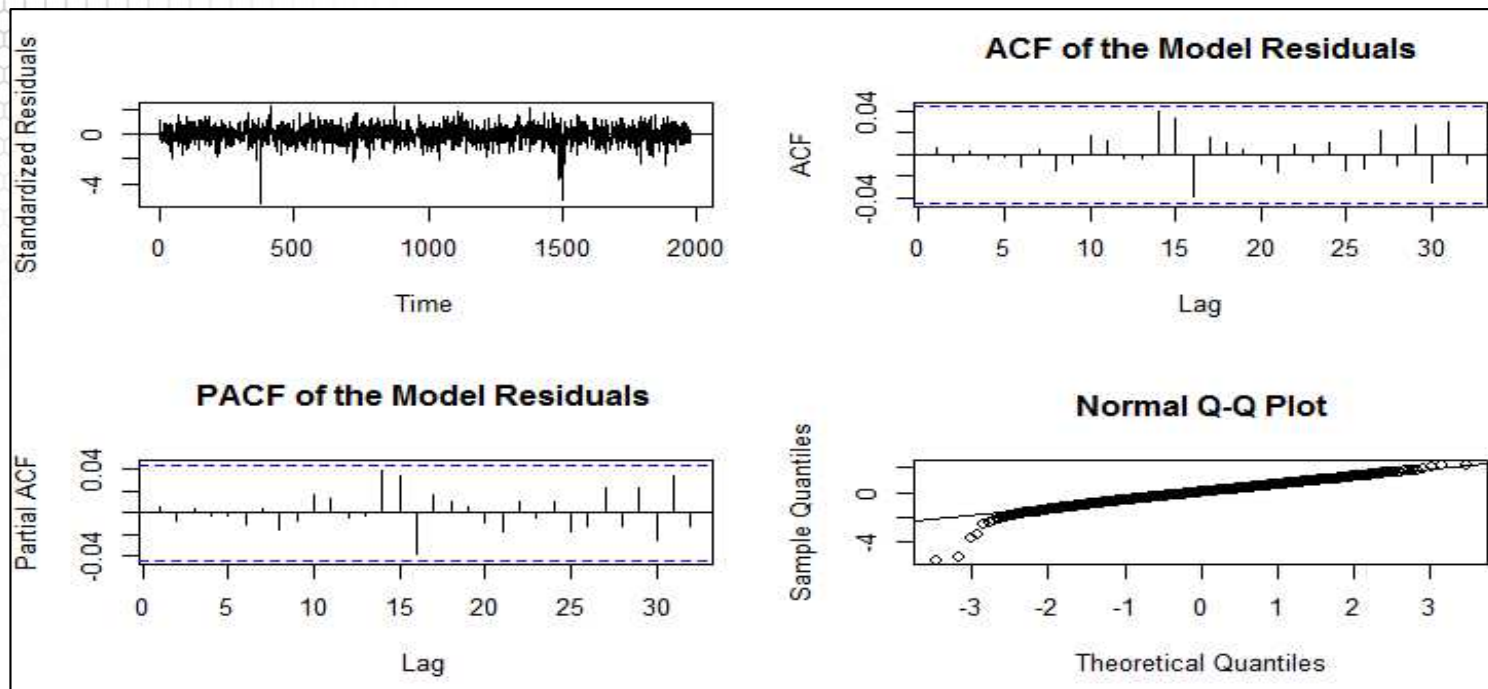
Which order to select?

```
aicv = as.vector(aic)
plot(aicv,ylab="AIC values")
indexp = rep(c(1:norder),norder)
indexq = rep(c(1:norder),each=norder)
indexaic = which(aicv == min(aicv))
porder = indexp[indexaic]-1
qorder = indexq[indexaic]-1
```

Final Model

```
final_model = arima(resid.process,order = c(porder,0,qorder), method='ML')
```


ARMA Model: Residual Analysis



Testing for Uncorrelated Residuals

Test for Uncorrelated Residuals for the final model

```
Box.test(final_model$resid, lag = (porder+qorder+1), type = "Box-Pierce", fitdf = (porder+qorder))
```

```
Box.test(final_model$resid, lag = (porder+qorder+1), type = "Ljung-Box", fitdf = (porder+qorder))
```

Test for Uncorrelated Residuals for the smaller model

```
Box.test(modarma$resid, lag = (porder+qorder+1), type = "Box-Pierce", fitdf = (porder+qorder))
```

```
Box.test(modarma$resid, lag = (porder+qorder+1), type = "Ljung-Box", fitdf = (porder+qorder))
```

Testing for Uncorrelated Residuals (cont'd)

```
> Box.test(final_model$resid, lag = (porder+qorder+1), type = "Box-Pierce", fitdf = (porder+qorder))
```

Box-Pierce test

```
data: final_model$resid  
X-squared = 1.9958, df = 1, p-value = 0.1577
```

```
> Box.test(final_model$resid, lag = (porder+qorder+1), type = "Ljung-Box", fitdf = (porder+qorder))
```

Box-Ljung test

```
data: final_model$resid  
X-squared = 2.0062, df = 1, p-value = 0.1567
```

```
> Box.test(modarma$resid, lag = (porder+qorder+1), type = "Box-Pierce", fitdf = (porder+qorder))
```

Box-Pierce test

```
data: modarma$resid  
X-squared = 2.2056, df = 1, p-value = 0.1375
```

```
> Box.test(modarma$resid, lag = (porder+qorder+1), type = "Ljung-Box", fitdf = (porder+qorder))
```

Box-Ljung test

```
data: modarma$resid  
X-squared = 2.2186, df = 1, p-value = 0.1364
```

Summary

