

Time Series Analysis

Modeling Heteroskedasticity

Nicoleta Serban, Ph.D.

Professor

Stewart School of Industrial and Systems Engineering

GARCH Model: Data Example

About This Lesson



PDC Energy, Inc (PDCE)

Summary:

- Crude oil and natural gas producer with headquartered in Denver, Colorado
- PDC's portfolio is comprised of the Wattenberg Field in Colorado, the Delaware Basin in West Texas and the Utica Shale in Ohio

Time Series Data:

- Daily stock price for more than 12 years of data starting with January 2007
- Largely dependent on the crude oil price



ARMA-GARCH Fit

Divide time series into training and testing

Predict August & September 1-16

```
pdcert2 = pdcert[-1]
```

```
n=length(pdcert2)
```

```
pdcert.test = pdcert2[3419:n]
```

```
pdcert.train = pdcert2[-c(3419:n)]
```

ARMA(4,4) & GARCH(1,1)

```
library(fGarch)
```

```
garchFit.ts = garchFit(~ arma(4,4)+ garch(1,1), data=pdcert.train, trace = FALSE)
```

```
fore_garch11 = predict(garchFit.ts, n.ahead = 32)
```



Prediction of 'garchFit' does not work properly
when considering joint arma+garch

ARMA-GARCH Fit: Different Implementation

ugrach from rugarch library (more computationally expensive)

ARMA(4,4) & GARCH(1,1)

```
library(rugarch)
```

```
spec = ugarchspec(variance.model=list(garchOrder=c(1,1)),  
                  mean.model=list(armaOrder=c(4,4), include.mean=T),  
                  distribution.model="std")
```

```
fit = ugarchfit(spec, pdcert.train, solver = 'hybrid')
```

```
fore = ugarchforecast(fit, n.ahead=32)
```



Prediction using 'ugarchfit' works
properly when considering joint models

ARMA+GARCH Order Selection

ARMA & GARCH – simultaneous fit and model selection for efficient estimators of the model parameters.


➡ ARMA(p, q)+GARCH(m, n) – computationally expensive to search over all combinations of $(p, q) \times (m, n)$ orders, hence apply a heuristic algorithm

- Step 1: Apply ARMA to model the time series and select orders \Rightarrow selected initial orders (p_0, q_0)
- Step 2: Apply ARMA(p_0, q_0)+GARCH(m, n) with varying m & n orders (consider only small values) \Rightarrow selected initial orders (m_0, n_0)
- Step 3: Apply ARMA(p, q)+GARCH(m_0, n_0) with varying p & q orders \Rightarrow selected initial orders (p_1, q_1)
- Step 4: Apply ARMA(p_1, q_1)+GARCH(m, n) with varying m & n orders (consider only small values) \Rightarrow selected initial orders (m_1, n_1)

Step 2: GARCH Order Selection

R function to be used across multiple combinations of (m,n) orders

```
test_modelAGG <- function(m,n){  
  spec <- ugarchspec(variance.model=list(garchOrder=c(m,n)),  
    mean.model=list(armaOrder=c(4,4),  
      include.mean=T),  
    distribution.model="std")  
  fit <- ugarchfit(spec, pdcert.train, solver = 'hybrid')  
  current.bic <- infocriteria(fit)[2]  
  df <- data.frame(m,n,current.bic)  
  names(df) <- c("m","n","BIC")  
  print(paste(m,n,current.bic,sep=" "))  
  return(df)  
}
```



Fix the orders for
modeling the
conditional mean:
ARMA(4,4)

Step 2: GARCH Order Selection (cont'd)

Consider all combinations of m & n between 0 to 2

```
ordersAGG = data.frame(Inf,Inf,Inf)
names/ordersAGG) <- c("m","n","BIC")
for (m in 0:2){
  for (n in 0:2){
    possibleError <- tryCatch(
      ordersAGG<-rbind(ordersAGG,test_modelAGG(m,n)),
      error=function(e) e
    )
    if(inherits(possibleError, "error")) next
  }
}
ordersAGG <- ordersAGG[order(-ordersAGG$BIC),]
```

Selected Orders for
modeling conditional
variance: GARCH(2,2)

Step 3: ARMA Order Selection Update

R function to be used across multiple combinations of (p,q) orders

```
test_modelAGA <- function(p,q){  
  spec = ugarchspec(variance.model=list(garchOrder=c(2,2)),  
                    mean.model=list(armaOrder=c(p,q),  
                                   include.mean=T),  
                    distribution.model="std")  
  fit = ugarchfit(spec, pdcert.train, solver = 'hybrid')  
  current.bic = infocriteria(fit)[2]  
  df = data.frame(p,q,current.bic)  
  names(df) <- c("p","q","BIC")  
  print(paste(p,q,current.bic,sep=" "))  
  return(df)  
}
```



Fix the orders for modeling
the conditional variance:
GARCH(2,2)

Step 3: ARMA Order Selection Update (cont'd)

Update the ARMA order

```
ordersAGA = data.frame(Inf,Inf,Inf)
names(ordersAGA) <- c("p","q","BIC")
for (p in 0:4){
  for (q in 0:4){
    possibleError <- tryCatch(
      ordersAGA<-rbind(ordersAGA,test_modelAGA(p,q)),
      error=function(e) e
    )
    if(inherits(possibleError, "error")) next
  }
}
ordersAGA <- ordersAGA[order(-ordersAGA$BIC),]
tail(ordersAGA)
```

Selected Orders for modeling the conditional mean:
ARMA(0,0)

Choose: AR=0,
MA=1 instead

Step 4: GARCH Order Selection Update

R function to be used across multiple combinations of (m,n) orders

```
test_modelAGG <- function(m,n){  
  spec = ugarchspec(variance.model=list(garchOrder=c(m,n)),  
    mean.model=list(armaOrder=c(0,1),  
    include.mean=T), distribution.model="std")  
  fit = ugarchfit(spec, pdcert.train, solver = 'hybrid')  
  current.bic = infocriteria(fit)[2]  
  df = data.frame(m,n,current.bic)  
  names(df) <- c("m","n","BIC")  
  print(paste(m,n,current.bic,sep=" "))  
  return(df)  
}
```

.....

Fix the orders for modeling the conditional mean: ARMA(0,1)

Selected Orders for modeling conditional variance: GARCH(1,1)

ARMA+GARCH Fit: Model Evaluation

Fit all three models and compare

```
spec.1 = ugarchspec(variance.model=list(garchOrder=c(2,2)),  
                    mean.model=list(armaOrder=c(4, 4),  
                                   include.mean=T), distribution.model="std")  
final.model.1 = ugarchfit(spec.1, pdcert.train, solver = 'hybrid')
```

.....

Compare Information Criteria

```
infocriteria(final.model.1)
```

```
infocriteria(final.model.2)
```

```
infocriteria(final.model.3)
```

```
## ARMA(4,4)+GARCH(2,2)
```

```
## ARMA(0,1)+GARCH(2,2)
```

```
## ARMA(0,1)+GARCH(1,1)
```

ARMA+GARCH Fit: Model Evaluation (cont'd)

> ## compare Information Criteria
> infocriteria(final.model.1)

Akaike	- 4.060253
Bayes	- 4.033321
Shibata	- 4.060291
Hannan-Quinn	- 4.050629

> infocriteria(final.model.2)

Akaike	- 4.055557
Bayes	- 4.041193
Shibata	- 4.055568
Hannan-Quinn	- 4.050424

> infocriteria(final.model.3)


Akaike	- 4.056774
Bayes	- 4.046001
Shibata	- 4.056780
Hannan-Quinn	- 4.052925



All models
perform similarly:
choose least
complex model


Residual Analysis

ACF of ARCH Residuals



White Noise

ACF of Squared ARCH Residuals



Not White Noise

Forecasting: Mean & Volatility

Prediction of the return time series and the volatility sigma

```
nfore = length(pdcert.test)
fore.series.1 = NULL
fore.sigma.1 = NULL
for(f in 1: nfore){
  ## Fit models
  data = pdcert.train
  if(f>2)
    data = c(pdcert.train,pdcert.test[1:(f-1)])
  final.model.1 = ugarchfit(spec.1, data, solver = 'hybrid')
  ## Forecast
  fore = ugarchforecast(final.model.1, n.ahead=1)
  fore.series.1 = c(fore.series.1, fore @forecast$seriesFor)
  fore.sigma.1 = c(fore.sigma.1, fore @forecast$sigmaFor)
}
```


Prediction Accuracy Comparison

```
> ### Mean Squared Prediction Error (MSPE)
> mean((fore.series.1 - pdcert.test)^2)
[1] 0.003071104
> mean((fore.series.2 - pdcert.test)^2)
[1] 0.003134157
> mean((fore.series.3 - pdcert.test)^2)
[1] 0.003133087
> ### Mean Absolute Prediction Error (MAE)
> mean(abs(fore.series.1 - pdcert.test))
[1] 0.03535444
> mean(abs(fore.series.2 - pdcert.test))
[1] 0.03671894
> mean(abs(fore.series.3 - pdcert.test))
[1] 0.03671133
```



Model 1 performs best
across all measures

```
> ### Mean Absolute Percentage Error (MAPE)
> mean(abs(fore.series.1 - pdcert.test)/abs(pdcert.test))
[1] 0.9554577
> mean(abs(fore.series.2 - pdcert.test)/abs(pdcert.test))
[1] 1.013034
> mean(abs(fore.series.3 - pdcert.test)/abs(pdcert.test))
[1] 1.012645
> ### Precision Measure (PM)
> sum((fore.series.1 - pdcert.test)^2)/sum((pdcert.test-mean(pdcert.test))^2)
[1] 0.9944229
> sum((fore.series.2 - pdcert.test)^2)/sum((pdcert.test-mean(pdcert.test))^2)
[1] 1.01484
> sum((fore.series.3 - pdcert.test)^2)/sum((pdcert.test-mean(pdcert.test))^2)
[1] 1.014493
```

Mean Prediction Comparison


Create a similar data structure for the forecasts

```
data.plot = pdcert.test  
names(data.plot)="Fore"
```

Compare observed time series with mean forecasts

```
plot(pdcert[c(n-90):n],type="l", ylim=c(ymin,ymax), xlab="Time", ylab="Return  
Price")  
data.plot$Fore=fore.series.1  
points(data.plot,lwd= 2, col="blue")  
data.plot$Fore=fore.series.2  
points(data.plot,lwd= 2, col="brown")  
data.plot$Fore=fore.series.3  
points(data.plot,lwd= 2, col="purple")
```

Mean Prediction Comparison




Variance Prediction Comparison

Compare squared observed time series with variance forecasts

```
plot(time.series^2,type="l", ylim=c(ymin,ymax), xlab="Time", ylab="Return Price")  
data.plot$Fore=fore.sigma.1^2  
points(data.plot,lwd= 2, col="blue")  
data.plot$Fore=fore.sigma.2^2  
points(data.plot,lwd= 2, col="brown")  
data.plot$Fore=fore.sigma.3^2  
points(data.plot,lwd= 2, col="purple")
```

Variance Prediction Comparison



Summary

