

# Introduction to R

---

# Road-map

- Data management
- Descriptive statistics
- Tests of association
- Linear regression and ANOVA
- Logistic and Cox proportional hazards regression

# Overview of R

- Free general-purpose statistical software
- Command line user interface
- Broad range of statistical functionality
- Excellent graphics generation capabilities
- *Extensive* additional functionality available through free downloadable packages
- User can write their own functions and scripts to further enhance and customize functionality

# R download

- R is free and versions for Windows or Mac (or Linux) can be downloaded from <http://www.r-project.org/> (choose your favorite mirror)

The screenshot shows the R Project for Statistical Computing website in a Windows Internet Explorer browser window. The browser's address bar displays <http://www.r-project.org/>. The website's main heading is "The R Project for Statistical Computing". On the left sidebar, the "Download, Packages" link is circled, with "CRAN" listed below it. The main content area features several statistical plots: a PCA plot titled "PCA 5 vars" showing variables like Fertility, Catholic, Examination, and Agriculture; a Clustering dendrogram titled "Clustering 4 groups"; and two histograms titled "Factor 1 [41%]" and "Factor 3 [19%]". The bottom of the page includes a "Getting Started:" section.

# A few R resources

- A few reference manuals at <http://www.r-project.org/>
- Short reference card available at <http://cran.r-project.org/doc/contrib/Short-refcard.pdf>

The screenshot shows the R Project for Statistical Computing website as it appeared in the early 2000s, viewed in a Windows Internet Explorer browser. The browser's address bar shows the URL <http://www.r-project.org/>. The website features the R logo on the left and a navigation menu with links such as "About R", "What is R?", "Contributors", "Screenshots", "What's new?", "Download, Packages", "CRAN", "R Project", "Foundation", "Members & Donors", "Mailing Lists", "Bug Tracking", "Developer Page", "Conferences", "Search", "Documentation", "Manuals", "FAQs", "The R Journal", and "What's new?". The "Manuals" link is circled in red. The main content area is titled "The R Project for Statistical Computing" and displays several statistical plots: a PCA plot of 5 variables (Fertility, Catholic, Agriculture, Examination, Education) with a correlation matrix, a scatter plot of Factor 1 [41%] vs Factor 3 [19%], a bar chart for Clustering 4 groups, and two histograms for Groups 1 and 2. The bottom of the page has a "Getting Started:" section.

# R Reference Card

by Tom Short, EPRI PEAC, tshort@epri-peac.com 2004-11-07

Granted to the public domain. See [www.Rpad.org](http://www.Rpad.org) for the source and latest version. Includes material from *R for Beginners* by Emmanuel Paradis (with permission).

## Getting help

Most R functions have online documentation.

`help(topic)` documentation on `topic`

`?topic id.`

`help.search("topic")` search the help system

`apropos("topic")` the names of all objects in the search list matching the regular expression "topic"

`help.start()` start the HTML version of help

`str(a)` display the internal "structure" of an R object

`summary(a)` gives a "summary" of `a`, usually a statistical summary but it is *generic* meaning it has different operations for different classes of `a`

`ls()` show objects in the search path; specify `pat="pat"` to search on a pattern

`ls.str()` `str()` for each variable in the search path

`dir()` show files in the current directory

`methods(a)` shows S3 methods of `a`

`methods(class=class(a))` lists all the methods to handle objects of class `a`

## Input and output

`load()` load the datasets written with `save`

`data(x)` loads specified data sets

`library(x)` load add-on packages

`read.table(file)` reads a file in table format and creates a data frame from it; the default separator `sep=""` is any whitespace; use `header=TRUE` to read the first line as a header of column names; use `as.is=TRUE` to prevent character vectors from being converted to factors; use `comment.char=""` to prevent "#" from being interpreted as a comment; use `skip=n` to skip `n` lines before reading data; see the help for options on row naming, NA treatment, and others

`read.csv("filename", header=TRUE)` id. but with defaults set for reading comma-delimited files

`read.delim("filename", header=TRUE)` id. but with defaults set for reading tab-delimited files

`read.fwf(file, widths, header=FALSE, sep=" ", as.is=FALSE)` read a table of fixed width formatted data into a 'data.frame'; `widths` is an integer vector, giving the widths of the fixed-width fields

`save(file, ...)` saves the specified objects (...) in the XDR platform-independent binary format

`save.image(file)` saves all objects

`cat(..., file="", sep=" ")` prints the arguments after coercing to character; `sep` is the character separator between arguments

`print(a, ...)` prints its arguments; *generic*, meaning it can have different methods for different objects

`format(x, ...)` format an R object for pretty printing

`write.table(x, file="", row.names=TRUE, col.names=TRUE, sep=" ")` prints `x` after converting to a data frame; if `quote` is `TRUE`,

character or factor columns are surrounded by quotes ("); `sep` is the field separator; `eol` is the end-of-line separator; `na` is the string for missing values; use `col.names=NA` to add a blank column header to get the column headers aligned correctly for spreadsheet input

`sink(file)` output to file, until `sink()`

Most of the I/O functions have a `file` argument. This can often be a character string naming a file or a connection. `file=""` means the standard input or output. Connections can include files, pipes, zipped files, and R variables.

On windows, the file connection can also be used with `description = "clipboard"`. To read a table copied from Excel, use

`x <- read.delim("clipboard")`

To write a table to the clipboard for Excel, use

`write.table(x, "clipboard", sep="\t", col.names=NA)`

For database interaction, see packages `RODBC`, `DBI`, `RMySQL`, `RpSQL`, and `ROracle`. See packages `XML`, `hdf5`, `netCDF` for reading other file formats.

## Data creation

`c(...)` generic function to combine arguments with the default forming a vector; with `recursive=TRUE` descends through lists combining all elements into one vector

`from:to` generates a sequence; ":" has operator priority; `1:4 + 1` is "2,3,4,5"

`seq(from,to)` generates a sequence `by=` specifies increment; `length=` specifies desired length

`seq(along=x)` generates `1, 2, ..., length(along)`; useful for for loops

`rep(x, times)` replicate `x` times; use `each=` to repeat "each" element of `x` each times; `rep(c(1,2,3), 2)` is `1 2 3 1 2 3`; `rep(c(1,2,3), each=2)` is `1 1 2 2 3 3`

`data.frame(...)` create a data frame of the named or unnamed arguments; `data.frame(v=1:4, ch=c("a", "B", "C", "d"), n=10)`; shorter vectors are recycled to the length of the longest

`list(...)` create a list of the named or unnamed arguments; `list(a=c(1,2), b="hi", c=3i)`

`array(x, dim=)` array with data `x`; specify dimensions like `dim=c(3, 4, 2)`; elements of `x` recycle if `x` is not long enough

`matrix(x, nrow=, ncol=)` matrix; elements of `x` recycle

`factor(x, levels=)` encodes a vector `x` as a factor

`gl(n, k, length=n*k, labels=1:n)` generate levels (factors) by specifying the pattern of their levels; `k` is the number of levels, and `n` is the number of replications

`expand.grid()` a data frame from all combinations of the supplied vectors or factors

`rbind(...)` combine arguments by rows for matrices, data frames, and others

`cbind(...)` id. by columns

## Slicing and extracting data

Indexing vectors

<code>x[n]</code>	$n^{th}$ element
<code>x[-n]</code>	all <i>but</i> the $n^{th}$ element
<code>x[1:n]</code>	first $n$ elements
<code>x[-(1:n)]</code>	elements from $n+1$ to the end
<code>x[c(1, 4, 2)]</code>	specific elements
<code>x["name"]</code>	element named "name"
<code>x[x &gt; 3]</code>	all elements greater than 3
<code>x[x &gt; 3 &amp; x &lt; 5]</code>	all elements between 3 and 5
<code>x[x %in% c("a", "and", "the")]</code>	elements in the given set

Indexing lists

<code>x[n]</code>	list with elements $n$
<code>x[[n]]</code>	$n^{th}$ element of the list
<code>x[["name"]]</code>	element of the list named "name"
<code>x\$name</code>	id.

Indexing matrices

<code>x[i, j]</code>	element at row $i$ , column $j$
<code>x[i, ]</code>	row $i$
<code>x[, j]</code>	column $j$
<code>x[, c(1, 3)]</code>	columns 1 and 3
<code>x["name", ]</code>	row named "name"

Indexing data frames (matrix indexing plus the following)

<code>x[["name"]]</code>	column named "name"
<code>x\$name</code>	id.

## Variable conversion

`as.array(x)`, `as.data.frame(x)`, `as.numeric(x)`,  
`as.logical(x)`, `as.complex(x)`, `as.character(x)`,  
... convert type; for a complete list, use `methods(as)`

## Variable information

`is.na(x)`, `is.null(x)`, `is.array(x)`, `is.data.frame(x)`,  
`is.numeric(x)`, `is.complex(x)`, `is.character(x)`,  
... test for type; for a complete list, use `methods(is)`

`length(x)` number of elements in `x`

`dim(x)` Retrieve or set the dimension of an object; `dim(x) <- c(3, 2)`

`dimnames(x)` Retrieve or set the dimension names of an object

`nrow(x)` number of rows; `NROW(x)` is the same but treats a vector as a one-row matrix

`ncol(x)` and `NCOL(x)` id. for columns

`class(x)` get or set the class of `x`; `class(x) <- "myclass"`

`unclass(x)` remove the class attribute of `x`

`attr(x, which)` get or set the attribute `which` of `x`

`attributes(obj)` get or set the list of attributes of `obj`

## Data selection and manipulation

`which.max(x)` returns the index of the greatest element of `x`

`which.min(x)` returns the index of the smallest element of `x`

`rev(x)` reverses the elements of `x`

`sort(x)` sorts the elements of `x` in increasing order; to sort in decreasing order: `rev(sort(x))`

`cut(x, breaks)` divides `x` into intervals (factors); `breaks` is the number of cut intervals or a vector of cut points

`match(x, y)` returns a vector of the same length than `x` with the elements of `x` which are in `y` (NA otherwise)

`which(x == a)` returns a vector of the indices of `x` if the comparison operation is true (TRUE), in this example the values of `i` for which `x[i] == a` (the argument of this function must be a variable of mode logical)

`choose(n, k)` computes the combinations of  $k$  events among  $n$  repetitions  
 $= n! / [(n-k)!k!]$

`na.omit(x)` suppresses the observations with missing data (NA) (suppresses the corresponding line if `x` is a matrix or a data frame)

`na.fail(x)` returns an error message if `x` contains at least one NA

**unique(x)** if **x** is a vector or a data frame, returns a similar object but with the duplicate elements suppressed

**table(x)** returns a table with the numbers of the different values of **x** (typically for integers or factors)

**subset(x, ...)** returns a selection of **x** with respect to criteria (...), typically comparisons: **x\$V1 < 10**; if **x** is a data frame, the option **select** gives the variables to be kept or dropped using a minus sign

**sample(x, size)** resample randomly and without replacement **size** elements in the vector **x**, the option **replace = TRUE** allows to resample with replacement

**prop.table(x, margin=)** table entries as fraction of marginal table

## Math

**sin, cos, tan, asin, acos, atan, atan2, log, log10, exp**

**max(x)** maximum of the elements of **x**

**min(x)** minimum of the elements of **x**

**range(x)** id. then **c(min(x), max(x))**

**sum(x)** sum of the elements of **x**

**diff(x)** lagged and iterated differences of vector **x**

**prod(x)** product of the elements of **x**

**mean(x)** mean of the elements of **x**

**median(x)** median of the elements of **x**

**quantile(x, probs=)** sample quantiles corresponding to the given probabilities (defaults to 0,.25,.5,.75,1)

**weighted.mean(x, w)** mean of **x** with weights **w**

**rank(x)** ranks of the elements of **x**

**var(x)** or **cov(x)** variance of the elements of **x** (calculated on  $n-1$ ); if **x** is a matrix or a data frame, the variance-covariance matrix is calculated

**sd(x)** standard deviation of **x**

**cor(x)** correlation matrix of **x** if it is a matrix or a data frame (1 if **x** is a vector)

**var(x, y)** or **cov(x, y)** covariance between **x** and **y**, or between the columns of **x** and those of **y** if they are matrices or data frames

**cor(x, y)** linear correlation between **x** and **y**, or correlation matrix if they are matrices or data frames

**round(x, n)** rounds the elements of **x** to **n** decimals

**log(x, base)** computes the logarithm of **x** with base **base**

**scale(x)** if **x** is a matrix, centers and reduces the data; to center only use the option **center=FALSE**, to reduce only **scale=FALSE** (by default **center=TRUE**, **scale=TRUE**)

**pmin(x, y, ...)** a vector which *i*th element is the minimum of **x[i]**, **y[i]**,...

**pmax(x, y, ...)** id. for the maximum

**cumsum(x)** a vector which *i*th element is the sum from **x[1]** to **x[i]**

**cumprod(x)** id. for the product

**cummin(x)** id. for the minimum

**cummax(x)** id. for the maximum

**union(x, y), intersect(x, y), setdiff(x, y), setequal(x, y), is.element(e1, set)** “set” functions

**Re(x)** real part of a complex number

**Im(x)** imaginary part

**Mod(x)** modulus; **abs(x)** is the same

**Arg(x)** angle in radians of the complex number

**Conj(x)** complex conjugate

**convolve(x, y)** compute the several kinds of convolutions of two sequences

**fft(x)** Fast Fourier Transform of an array

**mvfft(x)** FFT of each column of a matrix

**filter(x, filter)** applies linear filtering to a univariate time series or to each series separately of a multivariate time series

Many math functions have a logical parameter **na.rm=FALSE** to specify missing data (NA) removal.

## Matrices

**t(x)** transpose

**diag(x)** diagonal

**%%** matrix multiplication

**solve(a, b)** solves **a %% x = b** for **x**

**solve(a)** matrix inverse of **a**

**rowsum(x)** sum of rows for a matrix-like object; **rowSums(x)** is a faster version

**colsum(x), colSums(x)** id. for columns

**rowMeans(x)** fast version of row means

**colMeans(x)** id. for columns

## Advanced data processing

**apply(X, INDEX, FUN=)** a vector or array or list of values obtained by applying a function **FUN** to margins (**INDEX**) of **X**

**lapply(X, FUN)** apply **FUN** to each element of the list **X**

**tapply(X, INDEX, FUN=)** apply **FUN** to each cell of a ragged array given by **X** with indexes **INDEX**

**by(data, INDEX, FUN)** apply **FUN** to data frame **data** subsetted by **INDEX**

**merge(a, b)** merge two data frames by common columns or row names

**xtabs(a ~ b, data=x)** a contingency table from cross-classifying factors

**aggregate(x, by, FUN)** splits the data frame **x** into subsets, computes summary statistics for each, and returns the result in a convenient form; **by** is a list of grouping elements, each as long as the variables in **x**

**stack(x, ...)** transform data available as separate columns in a data frame or list into a single column

**unstack(x, ...)** inverse of **stack()**

**reshape(x, ...)** reshapes a data frame between ‘wide’ format with repeated measurements in separate columns of the same record and ‘long’ format with the repeated measurements in separate records; use (**direction=“wide”**) or (**direction=“long”**)

## Strings

**paste(...)** concatenate vectors after converting to character; **sep=** is the string to separate terms (a single space is the default); **collapse=** is an optional string to separate “collapsed” results

**substr(x, start, stop)** substrings in a character vector; can also assign, as **substr(x, start, stop) <- value**

**strsplit(x, split)** split **x** according to the substring **split**

**grep(pattern, x)** searches for matches to **pattern** within **x**; see **?regex**

**gsub(pattern, replacement, x)** replacement of matches determined by regular expression matching **sub()** is the same but only replaces the first occurrence.

**tolower(x)** convert to lowercase

**toupper(x)** convert to uppercase

**match(x, table)** a vector of the positions of first matches for the elements of **x** among **table**

**x %in% table** id. but returns a logical vector

**pmatch(x, table)** partial matches for the elements of **x** among **table**

**nchar(x)** number of characters

## Dates and Times

The class **Date** has dates without times. **POSIXct** has dates and times, including time zones. Comparisons (e.g. **>**), **seq()**, and **difftime()** are useful. **Date** also allows **+** and **-**. **?DateTimeClasses** gives more information. See also package **chron**.

**as.Date(s)** and **as.POSIXct(s)** convert to the respective class; **format(dt)** converts to a string representation. The default string format is “2001-02-21”. These accept a second argument to specify a format for conversion. Some common formats are:

**%a, %A** Abbreviated and full weekday name.

**%b, %B** Abbreviated and full month name.

**%d** Day of the month (01–31).

**%H** Hours (00–23).

**%I** Hours (01–12).

**%j** Day of year (001–366).

**%m** Month (01–12).

**%M** Minute (00–59).

**%p** AM/PM indicator.

**%S** Second as decimal number (00–61).

**%U** Week (00–53); the first Sunday as day 1 of week 1.

**%w** Weekday (0–6, Sunday is 0).

**%W** Week (00–53); the first Monday as day 1 of week 1.

**%y** Year without century (00–99). Don’t use.

**%Y** Year with century.

**%z** (output only.) Offset from Greenwich; **-0800** is 8 hours west of.

**%Z** (output only.) Time zone as a character string (empty if not available).

Where leading zeros are shown they will be used on output but are optional on input. See **?strftime**.

## Plotting

**plot(x)** plot of the values of **x** (on the **y**-axis) ordered on the **x**-axis

**plot(x, y)** bivariate plot of **x** (on the **x**-axis) and **y** (on the **y**-axis)

**hist(x)** histogram of the frequencies of **x**

**barplot(x)** histogram of the values of **x**; use **horiz=FALSE** for horizontal bars

**dotchart(x)** if **x** is a data frame, plots a Cleveland dot plot (stacked plots line-by-line and column-by-column)

**pie(x)** circular pie-chart

**boxplot(x)** “box-and-whiskers” plot

**sunflowerplot(x, y)** id. than **plot()** but the points with similar coordinates are drawn as flowers which petal number represents the number of points

**stripplot(x)** plot of the values of **x** on a line (an alternative to **boxplot()** for small sample sizes)

**coplot(x~y | z)** bivariate plot of **x** and **y** for each value or interval of values of **z**

**interaction.plot(f1, f2, y)** if **f1** and **f2** are factors, plots the means of **y** (on the **y**-axis) with respect to the values of **f1** (on the **x**-axis) and of **f2** (different curves); the option **fun** allows to choose the summary statistic of **y** (by default **fun=mean**)



**matplot(x,y)** bivariate plot of the first column of **x** vs. the first one of **y**, the second one of **x** vs. the second one of **y**, etc.

**fourfoldplot(x)** visualizes, with quarters of circles, the association between two dichotomous variables for different populations (**x** must be an array with `dim=c(2, 2, k)`, or a matrix with `dim=c(2, 2)` if  $k=1$ )

**assocplot(x)** Cohen–Friendly graph showing the deviations from independence of rows and columns in a two dimensional contingency table

**mosaicplot(x)** ‘mosaic’ graph of the residuals from a log-linear regression of a contingency table

**pairs(x)** if **x** is a matrix or a data frame, draws all possible bivariate plots between the columns of **x**

**plot.ts(x)** if **x** is an object of class “ts”, plot of **x** with respect to time, **x** may be multivariate but the series must have the same frequency and dates

**ts.plot(x)** id. but if **x** is multivariate the series may have different dates and must have the same frequency

**qqnorm(x)** quantiles of **x** with respect to the values expected under a normal law

**qqplot(x, y)** quantiles of **y** with respect to the quantiles of **x**

**contour(x, y, z)** contour plot (data are interpolated to draw the curves), **x** and **y** must be vectors and **z** must be a matrix so that `dim(z)=c(length(x), length(y))` (**x** and **y** may be omitted)

**filled.contour(x, y, z)** id. but the areas between the contours are coloured, and a legend of the colours is drawn as well

**image(x, y, z)** id. but with colours (actual data are plotted)

**persp(x, y, z)** id. but in perspective (actual data are plotted)

**stars(x)** if **x** is a matrix or a data frame, draws a graph with segments or a star where each row of **x** is represented by a star and the columns are the lengths of the segments

**symbols(x, y, ...)** draws, at the coordinates given by **x** and **y**, symbols (circles, squares, rectangles, stars, thermometres or “boxplots”) which sizes, colours ... are specified by supplementary arguments

**termplot(mod.obj)** plot of the (partial) effects of a regression model (**mod.obj**)

The following parameters are common to many plotting functions:

**add=FALSE** if TRUE superposes the plot on the previous one (if it exists)

**axes=TRUE** if FALSE does not draw the axes and the box

**type="p"** specifies the type of plot, “p”: points, “l”: lines, “b”: points connected by lines, “o”: id. but the lines are over the points, “h”: vertical lines, “s”: steps, the data are represented by the top of the vertical lines, “S”: id. but the data are represented by the bottom of the vertical lines

**xlim=, ylim=** specifies the lower and upper limits of the axes, for example with `xlim=c(1, 10)` or `xlim=range(x)`

**xlab=, ylab=** annotates the axes, must be variables of mode character

**main=** main title, must be a variable of mode character

**sub=** sub-title (written in a smaller font)

## Low-level plotting commands

**points(x, y)** adds points (the option `type=` can be used)

**lines(x, y)** id. but with lines

**text(x, y, labels, ...)** adds text given by **labels** at coordinates (x,y); a typical use is: `plot(x, y, type="n"); text(x, y, names)`

**mtext(text, side=3, line=0, ...)** adds text given by **text** in the margin specified by `side` (see `axis()` below); `line` specifies the line from the plotting area

**segments(x0, y0, x1, y1)** draws lines from points (x0,y0) to points (x1,y1)

**arrows(x0, y0, x1, y1, angle= 30, code=2)** id. with arrows at points (x0,y0) if `code=2`, at points (x1,y1) if `code=1`, or both if `code=3`; `angle` controls the angle from the shaft of the arrow to the edge of the arrow head

**abline(a,b)** draws a line of slope **b** and intercept **a**

**abline(h=y)** draws a horizontal line at ordinate **y**

**abline(v=x)** draws a vertical line at abscissa **x**

**abline(lm.obj)** draws the regression line given by **lm.obj**

**rect(x1, y1, x2, y2)** draws a rectangle which left, right, bottom, and top limits are **x1**, **x2**, **y1**, and **y2**, respectively

**polygon(x, y)** draws a polygon linking the points with coordinates given by **x** and **y**

**legend(x, y, legend)** adds the legend at the point (x,y) with the symbols given by **legend**

**title()** adds a title and optionally a sub-title

**axis(side, vect)** adds an axis at the bottom (`side=1`), on the left (2), at the top (3), or on the right (4); **vect** (optional) gives the abscissa (or ordinates) where tick-marks are drawn

**rug(x)** draws the data **x** on the *x*-axis as small vertical lines

**locator(n, type="n", ...)** returns the coordinates (x,y) after the user has clicked **n** times on the plot with the mouse; also draws symbols (`type="p"`) or lines (`type="l"`) with respect to optional graphic parameters (...); by default nothing is drawn (`type="n"`)

## Graphical parameters

These can be set globally with **par(...)**; many can be passed as parameters to plotting commands.

**adj** controls text justification (0 left-justified, 0.5 centred, 1 right-justified)

**bg** specifies the colour of the background (ex. : `bg="red"`, `bg="blue"`, ... the list of the 657 available colours is displayed with `colors()`)

**bty** controls the type of box drawn around the plot, allowed values are: “o”, “l”, “7”, “c”, “u” ou “j” (the box looks like the corresponding character); if `bty="n"` the box is not drawn

**cex** a value controlling the size of texts and symbols with respect to the default; the following parameters have the same control for numbers on the axes, `cex.axis`, the axis labels, `cex.lab`, the title, `cex.main`, and the sub-title, `cex.sub`

**col** controls the color of symbols and lines; use color names: “red”, “blue” see `colors()` or as “#RRGGBB”; see `rgb()`, `hsv()`, `gray()`, and `rainbow()`; as for **cex** there are: `col.axis`, `col.lab`, `col.main`, `col.sub`

**font** an integer which controls the style of text (1: normal, 2: italics, 3: bold, 4: bold italics); as for **cex** there are: `font.axis`, `font.lab`, `font.main`, `font.sub`

**las** an integer which controls the orientation of the axis labels (0: parallel to the axes, 1: horizontal, 2: perpendicular to the axes, 3: vertical)

**lty** controls the type of lines, can be an integer or string (1: “solid”, 2: “dashed”, 3: “dotted”, 4: “dotdash”, 5: “longdash”, 6: “twodash”, or a string of up to eight characters (between “0” and “9”) which specifies alternatively the length, in points or pixels, of the drawn elements and the blanks, for example `lty="44"` will have the same effect than `lty=2`

**lwd** a numeric which controls the width of lines, default 1

**mar** a vector of 4 numeric values which control the space between the axes and the border of the graph of the form `c(bottom, left, top, right)`, the default values are `c(5.1, 4.1, 4.1, 2.1)`

**mfcol** a vector of the form `c(nr,nc)` which partitions the graphic window as a matrix of **nr** lines and **nc** columns, the plots are then drawn in columns

**mfrow** id. but the plots are drawn by row

**pch** controls the type of symbol, either an integer between 1 and 25, or any single character within “”

1 ○ 2 △ 3 + 4 × 5 ◇ 6 ▽ 7 ☒ 8 \* 9 ⊕ 10 ⊕ 11 ☒ 12 ☒ 13 ☒ 14 ☒ 15 ■  
16 ● 17 ▲ 18 ◆ 19 ● 20 ● 21 ○ 22 □ 23 ◇ 24 △ 25 ▽ \* \* . . × × a a ? ?

**ps** an integer which controls the size in points of texts and symbols

**pty** a character which specifies the type of the plotting region, “s”: square, “m”: maximal

**tick** a value which specifies the length of tick-marks on the axes as a fraction of the smallest of the width or height of the plot; if `tick=1` a grid is drawn

**tol** a value which specifies the length of tick-marks on the axes as a fraction of the height of a line of text (by default `tol=0.5`)

**xaxt** if `xaxt="n"` the *x*-axis is set but not drawn (useful in conjunction with `axis(side=1, ...)`)

**yaxt** if `yaxt="n"` the *y*-axis is set but not drawn (useful in conjunction with `axis(side=2, ...)`)

## Lattice (Trellis) graphics

**xyplot(y~x)** bivariate plots (with many functionalities)

**barchart(y~x)** histogram of the values of **y** with respect to those of **x**

**dotplot(y~x)** Cleveland dot plot (stacked plots line-by-line and column-by-column)

**densityplot(~x)** density functions plot

**histogram(~x)** histogram of the frequencies of **x**

**bwplot(y~x)** “box-and-whiskers” plot

**qqmath(~x)** quantiles of **x** with respect to the values expected under a theoretical distribution

**stripplot(y~x)** single dimension plot, **x** must be numeric, **y** may be a factor

**qq(y~x)** quantiles to compare two distributions, **x** must be numeric, **y** may be numeric, character, or factor but must have two ‘levels’

**splo** matrix of bivariate plots

**parallel(~x)** parallel coordinates plot

**levelplot(z~x\*y | g1\*g2)** coloured plot of the values of **z** at the coordinates given by **x** and **y** (**x**, **y** and **z** are all of the same length)

**wireframe(z~x\*y | g1\*g2)** 3d surface plot

**cloud(z~x\*y | g1\*g2)** 3d scatter plot



In the normal Lattice formula, `y ~ x|g1*g2` has combinations of optional conditioning variables `g1` and `g2` plotted on separate panels. Lattice functions take many of the same arguments as base graphics plus also `data=` the data frame for the formula variables and `subset=` for subsetting. Use `panel=` to define a custom panel function (see `apropos("panel")` and `?llines`). Lattice functions return an object of class `trellis` and have to be printed to produce the graph. Use `print(xyplot(...))` inside functions where automatic printing doesn't work. Use `lattice.theme` and `lset` to change Lattice defaults.

## Optimization and model fitting

`optim(par, fn, method = c("Nelder-Mead", "BFGS", "CG", "L-BFGS-B", "SANN"))` general-purpose optimization; `par` is initial values, `fn` is function to optimize (normally minimize)

`nlm(f, p)` minimize function `f` using a Newton-type algorithm with starting values `p`

`lm(formula)` fit linear models; formula is typically of the form `response termA + termB + ...`; use `I(x*y) + I(x^2)` for terms made of nonlinear components

`glm(formula, family=)` fit generalized linear models, specified by giving a symbolic description of the linear predictor and a description of the error distribution; `family` is a description of the error distribution and link function to be used in the model; see `?family`

`nls(formula)` nonlinear least-squares estimates of the nonlinear model parameters

`approx(x, y=)` linearly interpolate given data points; `x` can be an xy plotting structure

`spline(x, y=)` cubic spline interpolation

`loess(formula)` fit a polynomial surface using local fitting

Many of the formula-based modeling functions have several common arguments: `data=` the data frame for the formula variables, `subset=` a subset of variables used in the fit, `na.action=` action for missing values: `"na.fail"`, `"na.omit"`, or a function. The following generics often apply to model fitting functions:

`predict(fit, ...)` predictions from `fit` based on input data

`df.residual(fit)` returns the number of residual degrees of freedom

`coef(fit)` returns the estimated coefficients (sometimes with their standard-errors)

`residuals(fit)` returns the residuals

`deviance(fit)` returns the deviance

`fitted(fit)` returns the fitted values

`logLik(fit)` computes the logarithm of the likelihood and the number of parameters

`AIC(fit)` computes the Akaike information criterion or AIC

## Statistics

`aov(formula)` analysis of variance model

`anova(fit, ...)` analysis of variance (or deviance) tables for one or more fitted model objects

`density(x)` kernel density estimates of `x`

`binom.test()`, `pairwise.t.test()`, `power.t.test()`, `prop.test()`, `t.test()`, ... use `help.search("test")`

## Distributions

`rnorm(n, mean=0, sd=1)` Gaussian (normal)

`rexp(n, rate=1)` exponential

`rgamma(n, shape, scale=1)` gamma

`rpois(n, lambda)` Poisson

`rweibull(n, shape, scale=1)` Weibull

`rcauchy(n, location=0, scale=1)` Cauchy

`rbeta(n, shape1, shape2)` beta

`rt(n, df)` 'Student' (*t*)

`rf(n, df1, df2)` Fisher-Snedecor (*F*) ( $\chi^2$ )

`rchisq(n, df)` Pearson

`rbinom(n, size, prob)` binomial

`rgeom(n, prob)` geometric

`rhyper(nn, m, n, k)` hypergeometric

`rlogis(n, location=0, scale=1)` logistic

`rlnorm(n, meanlog=0, sdlog=1)` lognormal

`rnbinom(n, size, prob)` negative binomial

`runif(n, min=0, max=1)` uniform

`rwilcox(nn, m, n)`, `rsignrank(nn, n)` Wilcoxon's statistics

All these functions can be used by replacing the letter `r` with `d`, `p` or `q` to get, respectively, the probability density (`dfunc(x, ...)`), the cumulative probability density (`pfunc(x, ...)`), and the value of quantile (`qfunc(p, ...)`), with  $0 < p < 1$ .

## Programming

`function( arglist ) expr` function definition

`return(value)`

`if(cond) expr`

`if(cond) cons.expr else alt.expr`

`for(var in seq) expr`

`while(cond) expr`

`repeat expr`

`break`

`next`

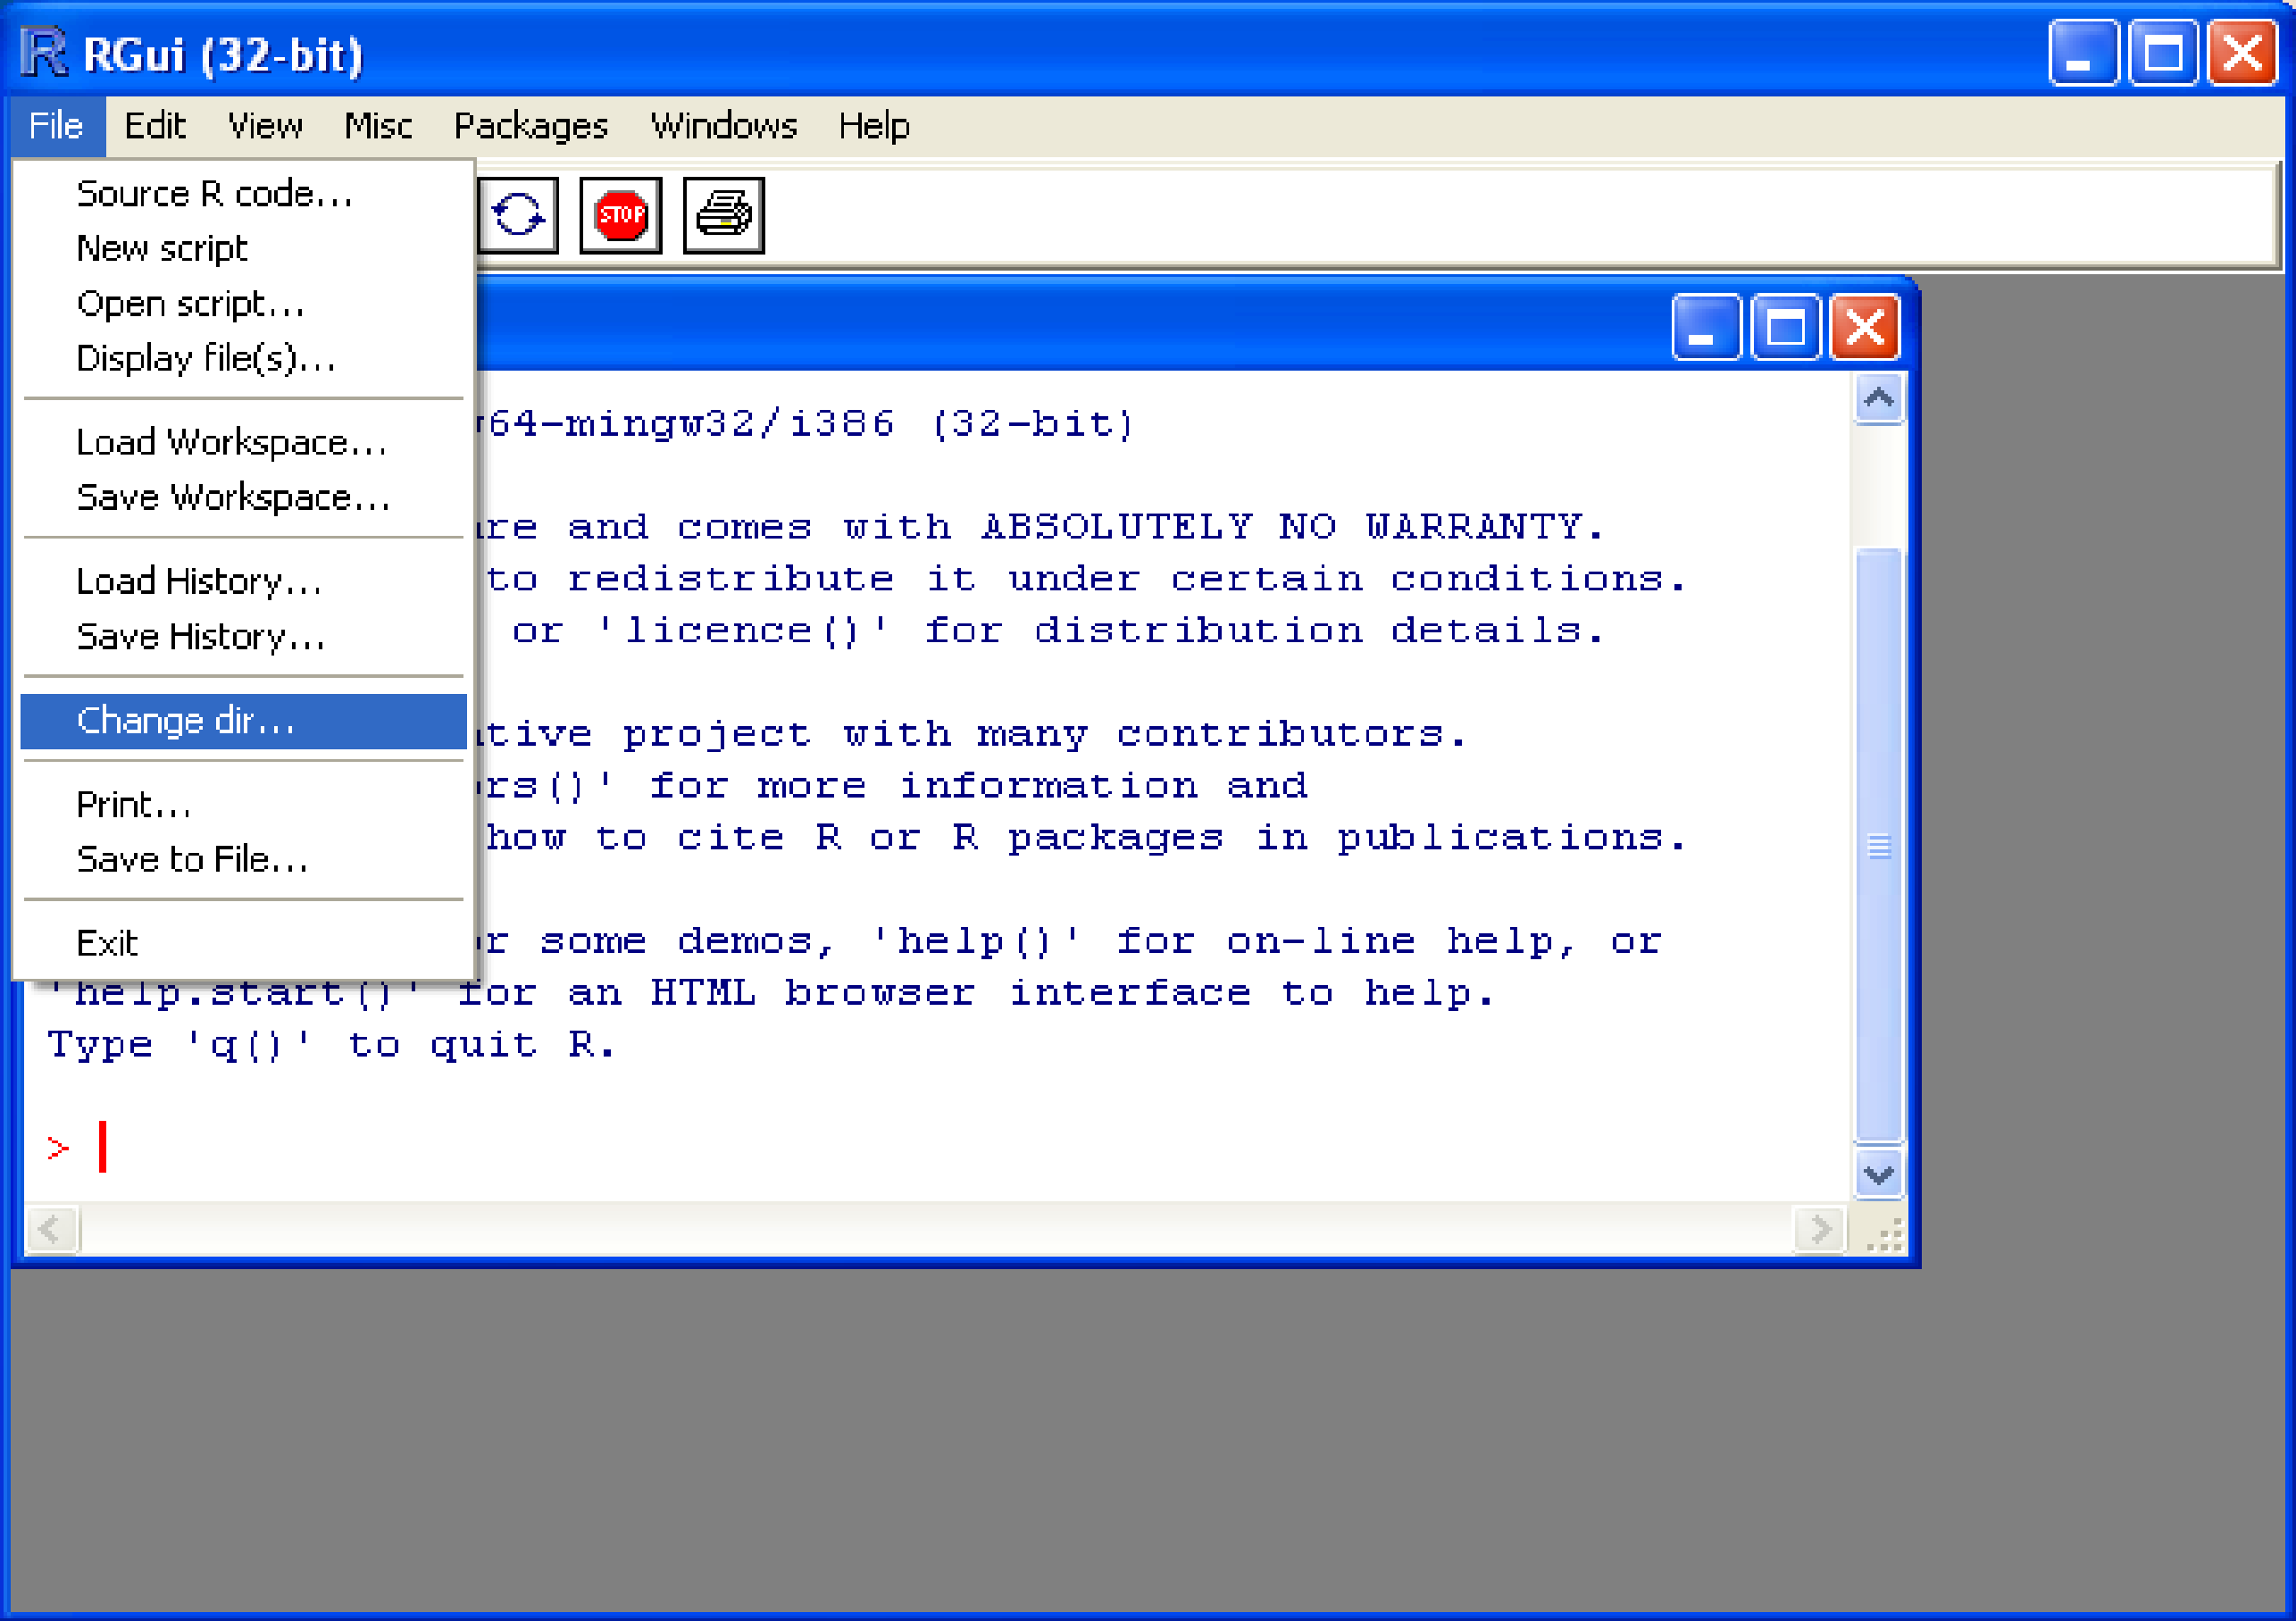
Use braces `{}` around statements

`ifelse(test, yes, no)` a value with the same shape as `test` filled with elements from either `yes` or `no`

`do.call(funname, args)` executes a function call from the name of the function and a list of arguments to be passed to it

# Changing the working directory in R

- Open R
- Change working directory
  - When loading data into R, the software will look in the working directory (unless directed otherwise)
  - Change to <relevant directory>.
  - “File”, “Change dir...” Then, browse to directory.



# Browse For Folder



Change working directory to:  
C:\Documents and Settings\gmsbah2\My Documents



Folder: My Documents

Make New Folder

OK

Cancel

NO WARRANTY.  
tain conditions.  
ation details.  
tributors.  
and  
in publications.  
on-line help, or  
to help.

# Installing packages in R

- Install packages
  - Particular types of functionality in R require the user to install packages
  - Install packages lme4 (generalized linear mixed models), geepack (generalized estimating equations models), and survival (survival analyses)
  - “Packages”, “Install package(s)...” Then, select mirror, and select package (packages have already been installed on these computers).






# Starting a new R script

- Typically, R commands are stored in a script which can be executed at the user's leisure or at a later date
- Open a new script in the R console. "File", "New script" Then, browse to directory.
- Click "Save as...", browse to selected location, and name the script.
- Don't forget ".R" or ".txt" after the file name (e.g. RworkshopScript.txt)

# Starting a new R script



# Miscellaneous items before importing the data

- **Naming rules/tips**
  - R is case sensitive (i.e. mydata and myData are different)
  - Use upper & lower case e.g. myData
- **Comments/help**
  - ?seq help(seq)
  - ?? "generalized estimating equations"
- **Logical tests (TRUE/FALSE)**

```
myNumbers=seq(0,18,3)
```

```
myNumbers==12
```

```
myNumbers<=5
```

```
myNumbers>10
```

# Miscellaneous items before importing the data

- **Objects**

- Scalars

- `myNumber=4`

- Vectors and matrices

- `myVector=c(1,2,3)`

- `myMatrix=matrix(seq(0,15,3),byrow=TRUE,ncol=2)`

- Dataframes store row and column names as well as column types (numeric, factor, etc.)

- Lists can store any type of R object

- **Indexing**

- Scalars have only one element (no indices)

- Individual entries in vectors, matrices, dataframes, and lists can be accessed using numeric indices or vectors of logicals (TRUE/FALSE)

- `myVector[2]`

- `myMatrix[3,]`

- `myMatrix[2,2]`

# Preparing data for R

- Data saved as .csv (comma delimited) is a good choice for loading into R.
- Other formats (e.g. SAS or SPSS) can *potentially* be loaded (foreign package)
- Open our datafiles “dataset.xls” and “data1.xls”  
Click “Save As”, select “CSV (Comma delimited)” from “Save as type:” box.
- Note that only selected sheet is saved
- Text fields with commas can cause problems

# Loading and examining data in R

#### Read data into R ####

```
myData=read.csv("dataset.csv",header=TRUE)
```

#### Display myData ####

```
myData
```

#### or ####

```
fix(myData)
```

#### Show attributes of myData (or any other object in R) ####

```
attributes(myData)
```

#### Show variable names myData ####

```
attributes(myData)$names
```



# Checking and assigning variable types in R

#### Examine the type of each variable ####

```
class(myData$patient_id)
class(myData$age)
class(myData$gender)
class(myData$diabetes)
class(myData$pain)
class(myData$health_status)
class(myData$mortality_12months)
class(myData$pain2)
```

#### Assign proper variable types to each variable ####

```
myData$patient_id=as.factor(myData$patient_id)
myData$age=as.numeric(myData$age)
myData$gender=as.factor(myData$gender)
myData$diabetes=as.factor(myData$diabetes)
myData$pain=as.numeric(myData$pain)
myData$health_status=as.factor(myData$health_status)
myData$mortality_12months=as.factor(myData$mortality_12months)
myData$pain2=as.numeric(myData$pain2)
```

# Summarizing variables in R

#### Summaries of variables ####

```
summary(myData$age)
summary(myData$gender)
summary(myData$diabetes)
summary(myData$pain)
summary(myData$health_status)
summary(myData$mortality_12months)
summary(myData$pain2)
```

```
mean(myData$age,na.rm=TRUE)
sd(myData$age,na.rm=TRUE)
```

```
median(myData$pain)
min(myData$pain)
max(myData$pain)
IQR(myData$pain)
```

#### A few more detailed summaries ####

#### Summary of women's ages ####

```
summary(myData$age[myData$gender=="0"])
```

#### Cross-tabulation of gender and diabetes  
####

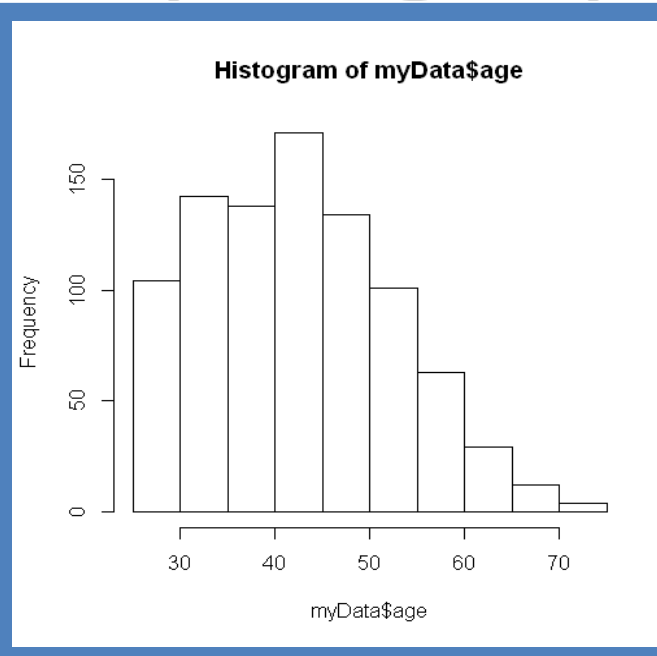
```
table(myData$gender,myData$diabetes)
```

# Generating simple graphs in R

#### A few graphical summaries ####


#### Histogram of age ####

```
hist(myData$age)
```



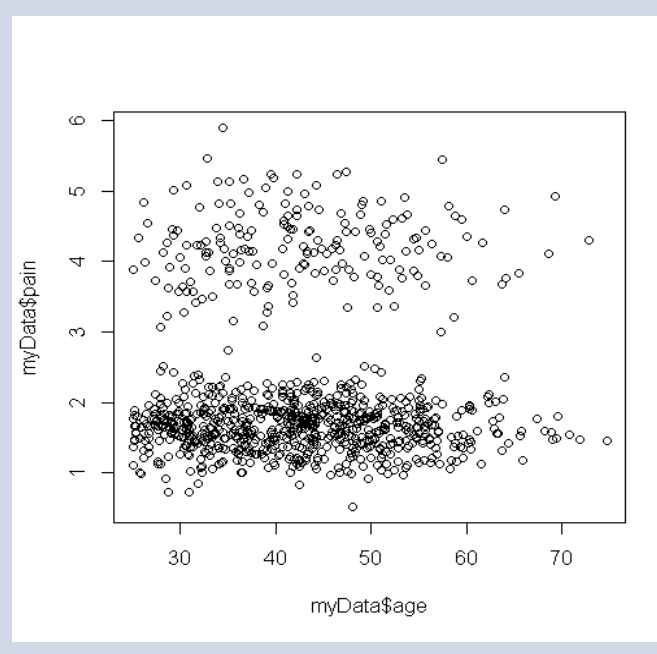
#### Barchart of health status ####

```
plot(myData$health_status)
```




#### Scatterplot of pain by age ####

```
plot(myData$age,myData$pain)
```



#### Boxplots of pain by health status ####

```
plot(myData$health_status,myData$pain)
```



- To save, highlight plot window in R console, “File”, “Save as”, and select format.

# A more complete plot in R

```
#### A more complete plot ####
```

```
dev.new(height=3,width=4)
```

```
par(mai=c(0.3,0.5,0.3,0.05),cex=0.8)
```

```
plot(myData$diabetes,myData$pain,  
     axes=FALSE,xlab="",ylab="",  
     main="Pain Score Vs. Diabetes")
```

```
box()
```

```
axis(1,at=c(1,2),labels=c("No Diabetes","Diabetes"),  
     padj=-0.5)
```

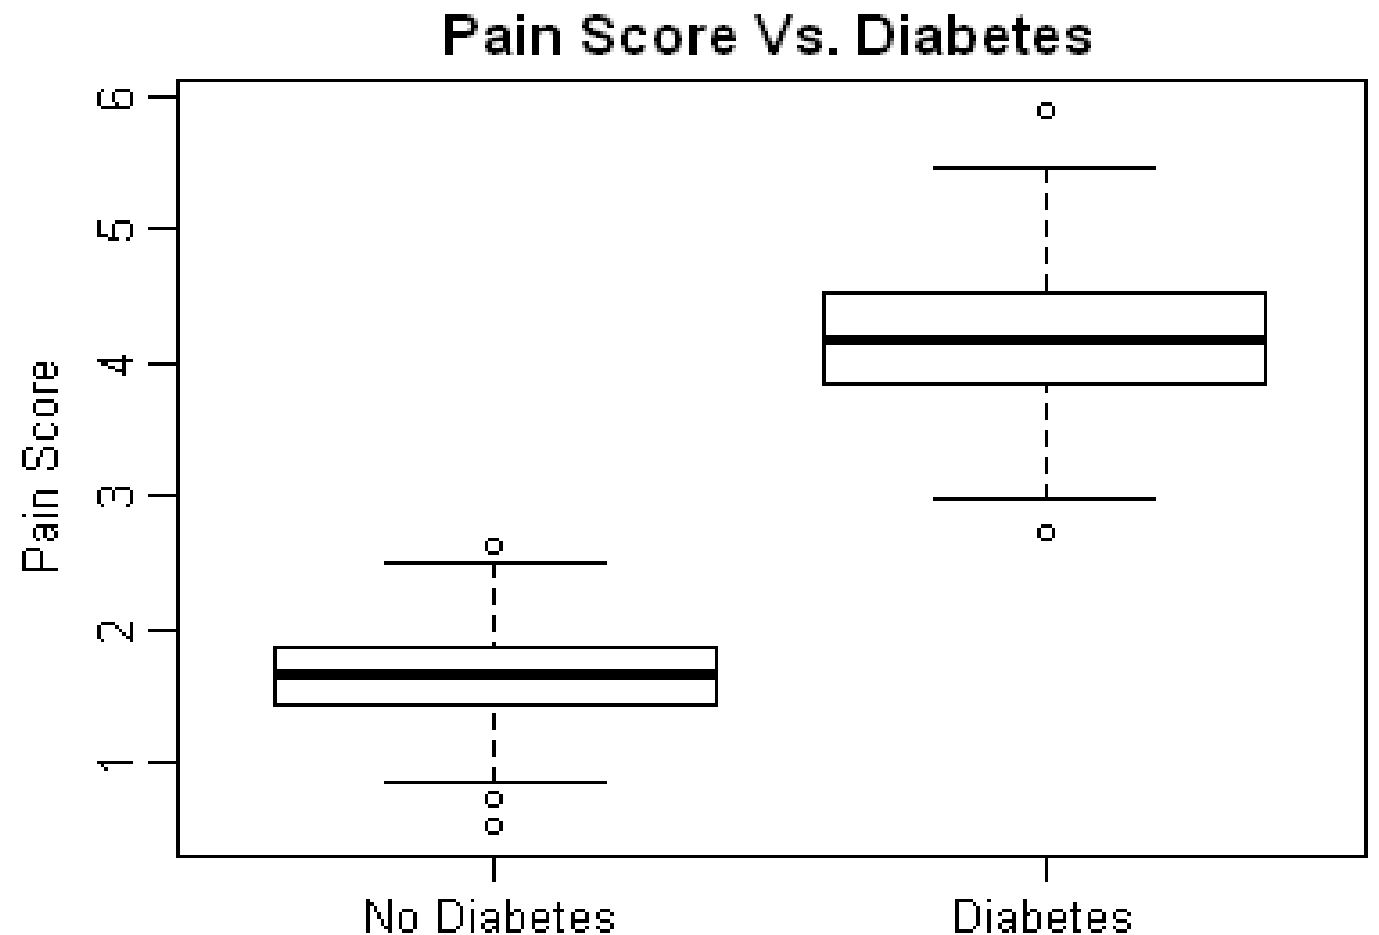
```
axis(2,padj=0.5)
```

```
title(ylab="Pain Score",line=2)
```

```
#### Check documentation for more options ####
```

```
?par
```

```
?plot
```



# Creating new variables in R

- Suppose we want a log(pain) variable and to categorize age into  $\leq 35$ , 36-45, 46-55,  $> 55$  years

#### Create logPain variable ####

```
logPain=log(myData$pain)
```

#### Create categorized age variable ####

```
ageCats=sapply(myData$age,function(x){  
  ifelse(x<=35,return("<=35"),  
    ifelse(x<=45,return("36-45"),  
      ifelse(x<=55,return("46-  
55"),return(">55"))))  
})
```

#### alternatively ####

```
ageCatsV2=cut(myData$age,breaks=c(0,35,45,5  
5,Inf))
```

#### how do they look? ####

```
cbind(ageCats,ageCatsV2)
```

#### use as.character(),

#### so that factor labels are displayed

#### instead of factor numbers

```
cbind(ageCats,as.character(ageCatsV2))
```

# Univariate tests of association for independent observations

- Comparing central tendency (e.g. means) in 2 or more independent groups (t-test, Wilcoxon rank sum test, ANOVA, Kruskal-Wallis)
- Comparing proportions in 2 or more groups (chi-squared test, Fisher's exact test)
- Tests of correlation (Pearson, Spearman)



# Comparing central tendency in 2 or more independent groups

- Consider comparing the mean ages of diabetics and non-diabetics

#### t-test of mean age by diabetes status  
#### (two groups) ####

```
t.test(myData$age~myData$diabetes)
```

#### Wilcoxon rank sum test of whether ages  
#### tend to be larger to smaller by  
#### diabetes status  
#### (two groups) ####

```
wilcox.test(myData$age~myData$diabetes)
```

- Consider comparing the mean pain scores across health status

#### ANOVA of pain by health status  
#### (two or more groups) ####

```
summary(lm(myData$pain~myData$health_status))
```

#### Kruskal-Wallis test of whether pain  
#### scores tend to be larger to smaller by  
#### health status  
#### (two or more groups) ####

```
kruskal.test(myData$pain~myData$health_status)
```

# Comparing proportions in 2 or more independent groups

- Consider comparing the proportion of mortalities at 12 months among diabetics and non-diabetics

```
#### Chi-squared test of mortality by diabetes ####
```

```
chisq.test(myData$mortality_12month,myData$diabetes)
```

```
#### Fisher's exact test of mortality by diabetes ####
```

```
fisher.test(myData$mortality_12month,myData$diabetes)
```

```
#### alternatively ####
```

```
myTable=table(myData$mortality_12month,myData$diabetes)
```

```
chisq.test(myTable)
```

```
fisher.test(myTable)
```

# Tests of correlation

- Consider testing for an association between age and pain score

#### test for correlation between age  
#### and pain score ####

```
cor.test(myData$age,myData$pain)
```

#### if we think the relationship may be  
#### non-linear ####

```
cor.test(myData$age,myData$pain,  
         method="spearman")
```

#### alternatively, linear regression ####

#### Wald test p-values ####

```
summary(lm(myData$age~myData$pain))
```

#### or ####

```
summary(lm(myData$pain~myData$age))
```

#### or ####

```
summary(lm(logPain~myData$age))
```

# Univariate tests of association for dependent observations

- Comparing central tendency (e.g. means) in 2 or more dependent groups (paired t-test, Wilcoxon signed rank test, linear mixed or GEE model)
- Comparing proportions in 2 or more dependent groups (McNemar's test, logistic mixed or GEE model)

# Comparing 2 or more dependent groups

- Consider a cross-over study where each subject receives both treatment and placebo (in a random order), a subject's response is considered successful if it exceeds 1

#### Generate some fake data!!!! ####

```
mySampleSize=100
```

```
personIndicator=rep(seq(1,mySampleSize,1),rep(2,mySampleSize))
```

```
personEffect=rnorm(n=mySampleSize,mean=0,sd=1)
personEffect=rep(personEffect,rep(2,mySampleSize))
```

```
treatmentIndicator=c(0,1)
treatmentIndicator=rep(treatmentIndicator,mySampleSize)
```

```
treatmentEffect=1.234*treatmentIndicator
```

```
myResponse=treatmentEffect+personEffect+rnorm(2*mySampleSize)
```

```
mySuccess=(myResponse>=1)
```

```
cbind(personIndicator,treatmentIndicator,myResponse,mySuccess)
```

# Comparing central tendency in 2 or more dependent groups

- Consider comparing the response in the treatment and placebo groups (need to account for fact that observations on the same subject are dependent)

#### paired t.test for treatment effect ####

```
t.test(myResponse[treatmentIndicator=="1"],  
      myResponse[treatmentIndicator=="0"],  
      paired=TRUE)
```

#### signed rank test for treatment effect ####

```
wilcox.test(myResponse[treatmentIndicator=="1"],  
            myResponse[treatmentIndicator=="0"],  
            paired=TRUE)
```

#### alternatively ####

```
library(geepack)  
library(lme4)
```

#### linear GEE model ####

```
summary(geeglm(myResponse~treatmentIndicator,  
              id=personIndicator))
```

#### linear mixed effects model ####

```
summary(lmer(myResponse~treatmentIndicator+  
            (1|personIndicator)))
```



# Comparing proportions in 2 or more dependent groups

- Consider comparing success probability in treatment and placebo groups using McNemar's test

#### McNemar's test for comparing probability of success

#### in treatment and control groups ####

```
table(mySuccess[treatmentIndicator=="0"],  
      mySuccess[treatmentIndicator=="1"])
```

```
mcnemar.test(mySuccess[treatmentIndicator=="0"],  
             mySuccess[treatmentIndicator=="1"],correct=FALSE)
```

#### alternatively ####

```
discordant=c(table(mySuccess[treatmentIndicator=="0"],  
                  mySuccess[treatmentIndicator=="1"])[1,2],  
             table(mySuccess[treatmentIndicator=="0"],  
                  mySuccess[treatmentIndicator=="1"])[2,1])
```

```
chisq.test(discordant,correct=FALSE)
```

#### or ####

```
2*(1-pbinom(discordant[1],sum(discordant),1/2))
```

#### alternatively ####

#### logistic GEE model ####

```
summary(geeglm(mySuccess~treatmentIndicator,  
              id=personIndicator,  
              family="binomial"))
```

#### logistic mixed effects model ####

```
summary(lmer(mySuccess~treatmentIndicator+  
            (1|personIndicator),  
            family="binomial"))
```

# Multivariate Analyses

```
#### First, some fake data! ####
```

```
myTreatment=rep(c("Trt 1","Trt 2"),c(50,50))
```

```
myConfounder=c(rnorm(50),rnorm(50)+5)
```

```
myResponse=c(myConfounder[1:50]-rnorm(50),  
             myConfounder[51:100]-rnorm(50)-7)
```

```
myFailure=(myResponse<0)
```

```
#### Examine relationship between
```

```
#### myResponse and myTreatment ####
```

```
summary(lm(myResponse~myTreatment))
```

```
#### Is the relationship modified
```

```
#### when we adjust by myConfounder? ####
```

```
summary(lm(myResponse~myConfounder+myTreatment))
```

```
#### Examine relationship between
```

```
#### myFailure and myTreatment ####
```

```
#### Logistic Regression ####
```

```
summary(glm(myFailure~myTreatment,  
            family="binomial"))
```

```
#### Is the relationship modified
```

```
#### when we adjust by myConfounder? ####
```

```
summary(glm(myFailure~myConfounder+myTreatment,  
            family="binomial"))
```

# Multivariate Analyses (continued)

#### And make a plot ####

```
dev.new(height=6,width=8)
par(mai=c(0.5,0.5,0.05,0.05))
```

```
plot(x=NULL,y=NULL,xlim=c(-2,8),ylim=c(-5,4),axes=FALSE)
box()
```

```
points(x=myConfounder[myTreatment=="Trt 1"],
       y=myResponse[myTreatment=="Trt 1"],
       pch="x")
points(x=myConfounder[myTreatment=="Trt 2"],
       y=myResponse[myTreatment=="Trt 2"],
       pch=21)
```

```
abline(lm(myResponse~myConfounder)$coef)
```


```
myCoefs=lm(myResponse~myTreatment+myConfounder)$coef
abline(myCoefs[1],myCoefs[3],lty="dotted")
abline(myCoefs[1]+myCoefs[2],myCoefs[3],lty="dotted")
```

```
title(xlab="Confounder",line=1.25)
```

```
title(ylab="Response",line=1.5)
```

```
legend(-2,4,legend=c("Trt 1","Trt 2"),pch=c("x","o"))
```

```
axis(1,at=seq(-2,8,1),padj=-1)
```



# Enrichment Work

1. Load the dataset “data1.csv” into R using the `read.csv()` command (don’t forget the `header=TRUE` option).

The dataset contains group indicators, log white blood cell counts, remission time subject to censoring by loss-to-follow up, and an event indicator.

2. Ensure that variables are of the appropriate type using the `as.factor()` and `as.numeric()` commands.

*The event indicator should be numeric for our purposes.*

# Enrichment Work (continued)

3. Make a box-plot of logWBC vs. group using the plot() command. Adjust the margins and label the axes by replacing the ???s in the below code.

```
par(mai=c(???,???,???,???),cex=???)  
plot(group,logWBC,axes=FALSE)  
box()  
axis(1,padj=-0.5)  
axis(2,padj=0.5)  
title(xlab=???,line=1.5)  
title(ylab=???,line=1.75)
```

# Enrichment Work (continued)

4. Load the survival library using the command `library(survival)`.
5. Create a survival response (subject to non-informative right censoring) using the `Surv()` command.

*The ingredients for this will be remission time and event indicator.*

# Enrichment Work (continued)

6. Plot Kaplan-Meier survival function estimates for each group by replacing the ???s in the below code.

```
par(mai=c(0.5,0.5,0.05,0.05),cex=0.8)
plot(survfit(???~???),lty=c("solid","dotted"),axes=FALSE)
box()
axis(1,adj=???)
axis(2,adj=???)
title(xlab="Remission Time",line=???)
title(ylab="Survival Probability",line=???)
legend(x="topright",
      legend=c("Group 0","Group 1"),
      lty=c("solid","dotted"))
```

*The survival response should go to the left of the “~” and the group indicator should go to the right of the “~”.*

7. Annotate the plot with the p-value from the log-rank test by replacing the ???s in the below code.

```
summary(coxph(mySurv~group))
text(x=35,y=0,labels="Log-Rank p ???",pos=2)
```

# Enrichment Work (continued)

8. Compare a Cox proportional hazards model adjusting for log WBC in addition to group to a model with group alone by updating the code below.

```
fit1=coxph(???~???)  
fit2=coxph(???~???+???)  
anova(fit1,fit2)
```

*The dependent variable goes on the left side of the “~” sign and predictors go on the right hand side of the “~”, separated by “+” if there are multiple predictors.*



# Enrichment Work (continued)

9. Extract hazard ratio estimates and 95% confidence intervals from fit2 using the summary() command.
10. How should these hazard ratio estimates and confidence intervals be interpreted?